

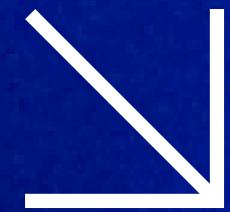
Optimizing Autonomous Agents in

# vacuum<sup>o</sup> World

Group 1

2026  
Artificial  
Intelligence

# Table of Contents



#1

Introduction & Problem Specification

#4

Informed Search & Heuristics

#2

Search Problem Formalization

#5

Experimental Results & Comparative Analysis

#3

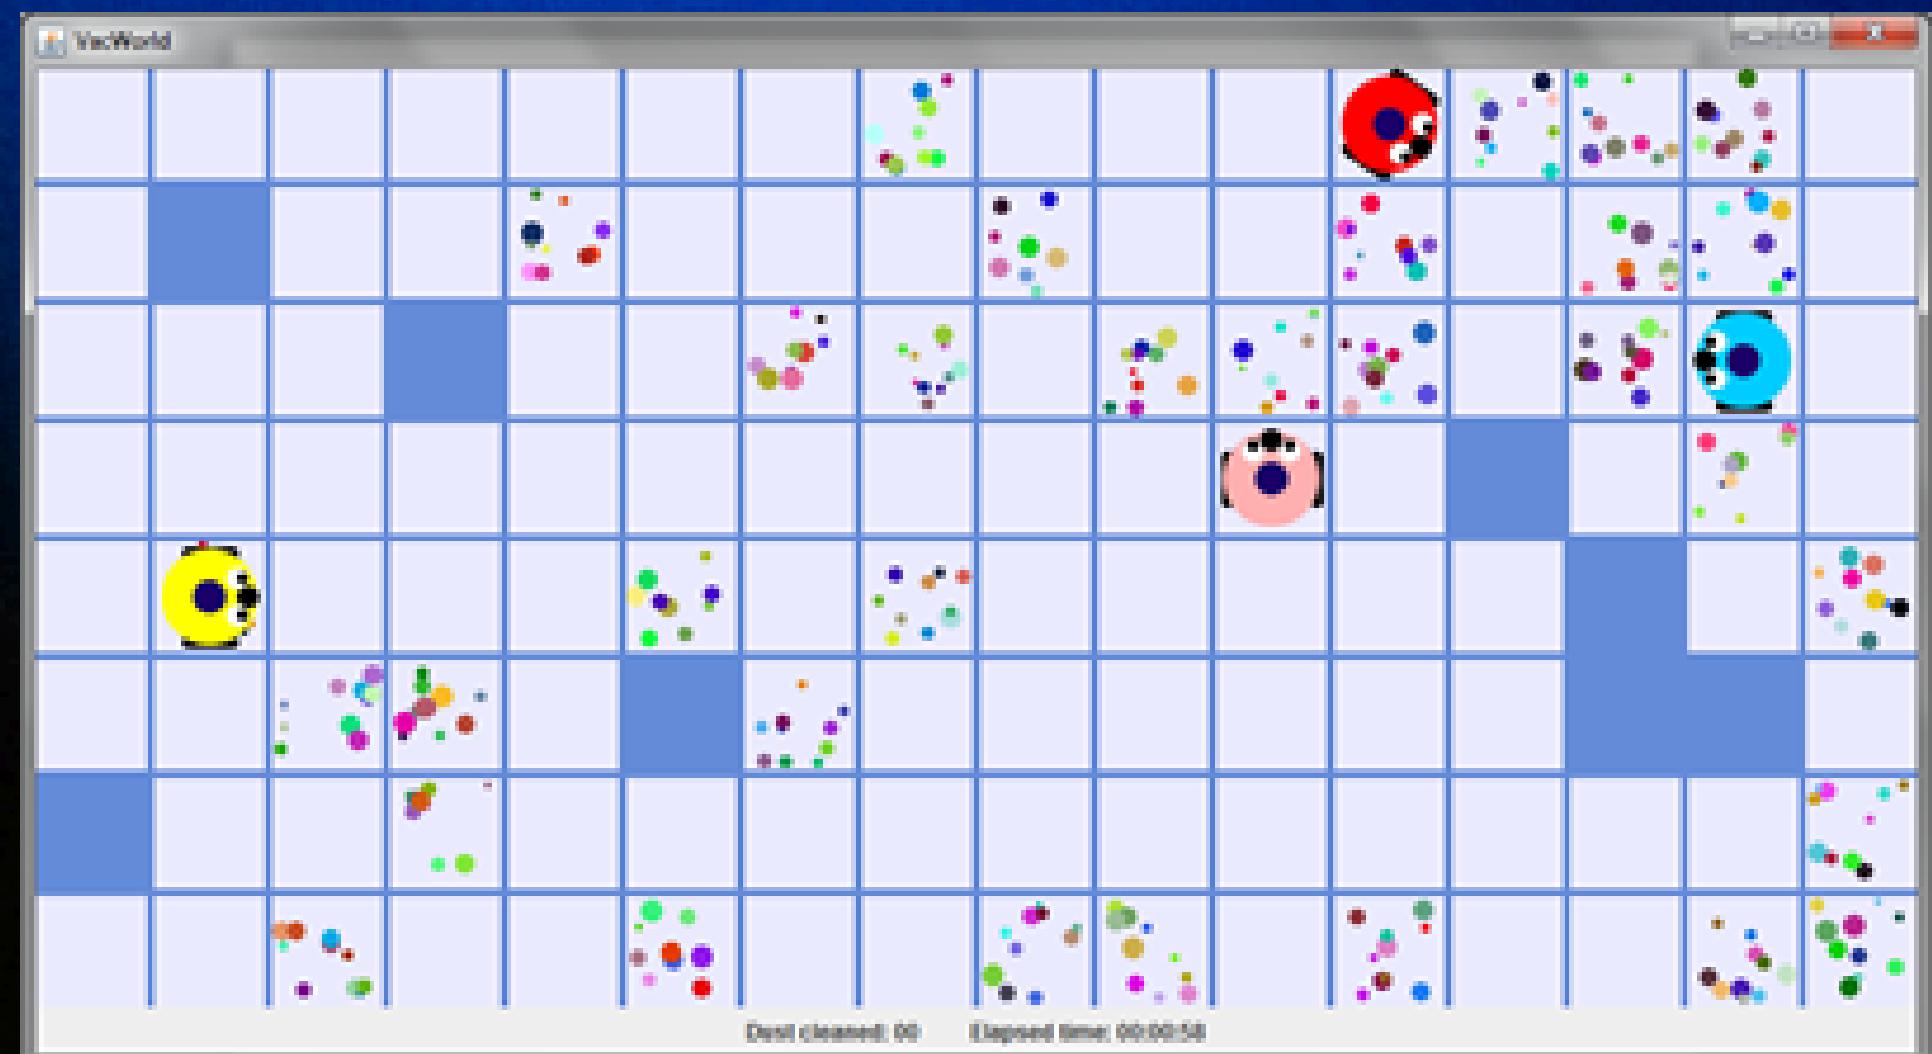
Search Tree & Uninformed Strategies

#6

Conclusion & Future Work



# Why Vacuum World



A classic AI benchmark for rational agents.

Demonstrates core concepts:  
State space, Search trees, and  
Heuristics.

## Goal

- » Autonomous cleaning with efficiency (path length vs. time).

## Challenges

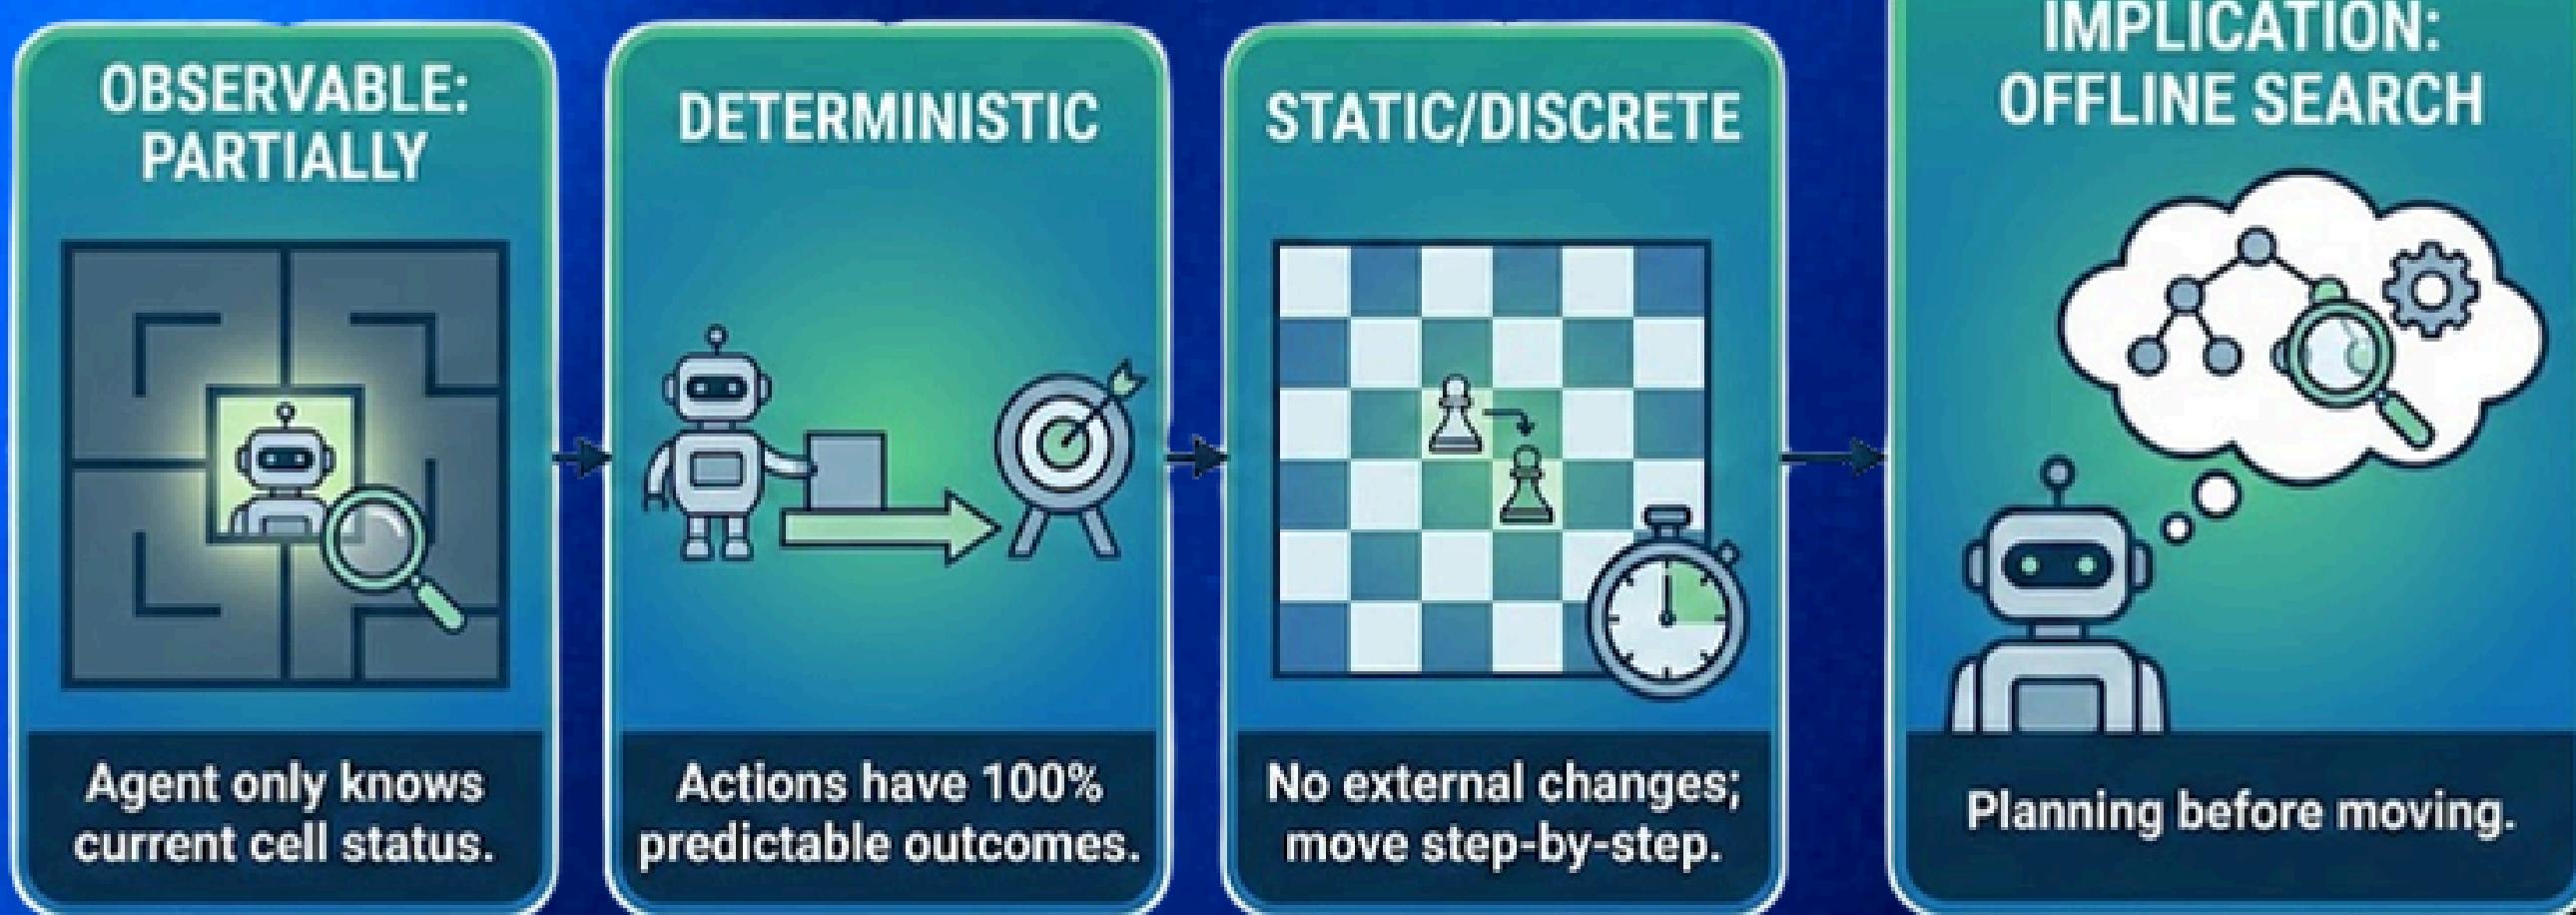
- » Exponential state growth as grid size ( $n \times n$ ) increases.



## Agent Specification: PEAS Framework



# Environment Properties

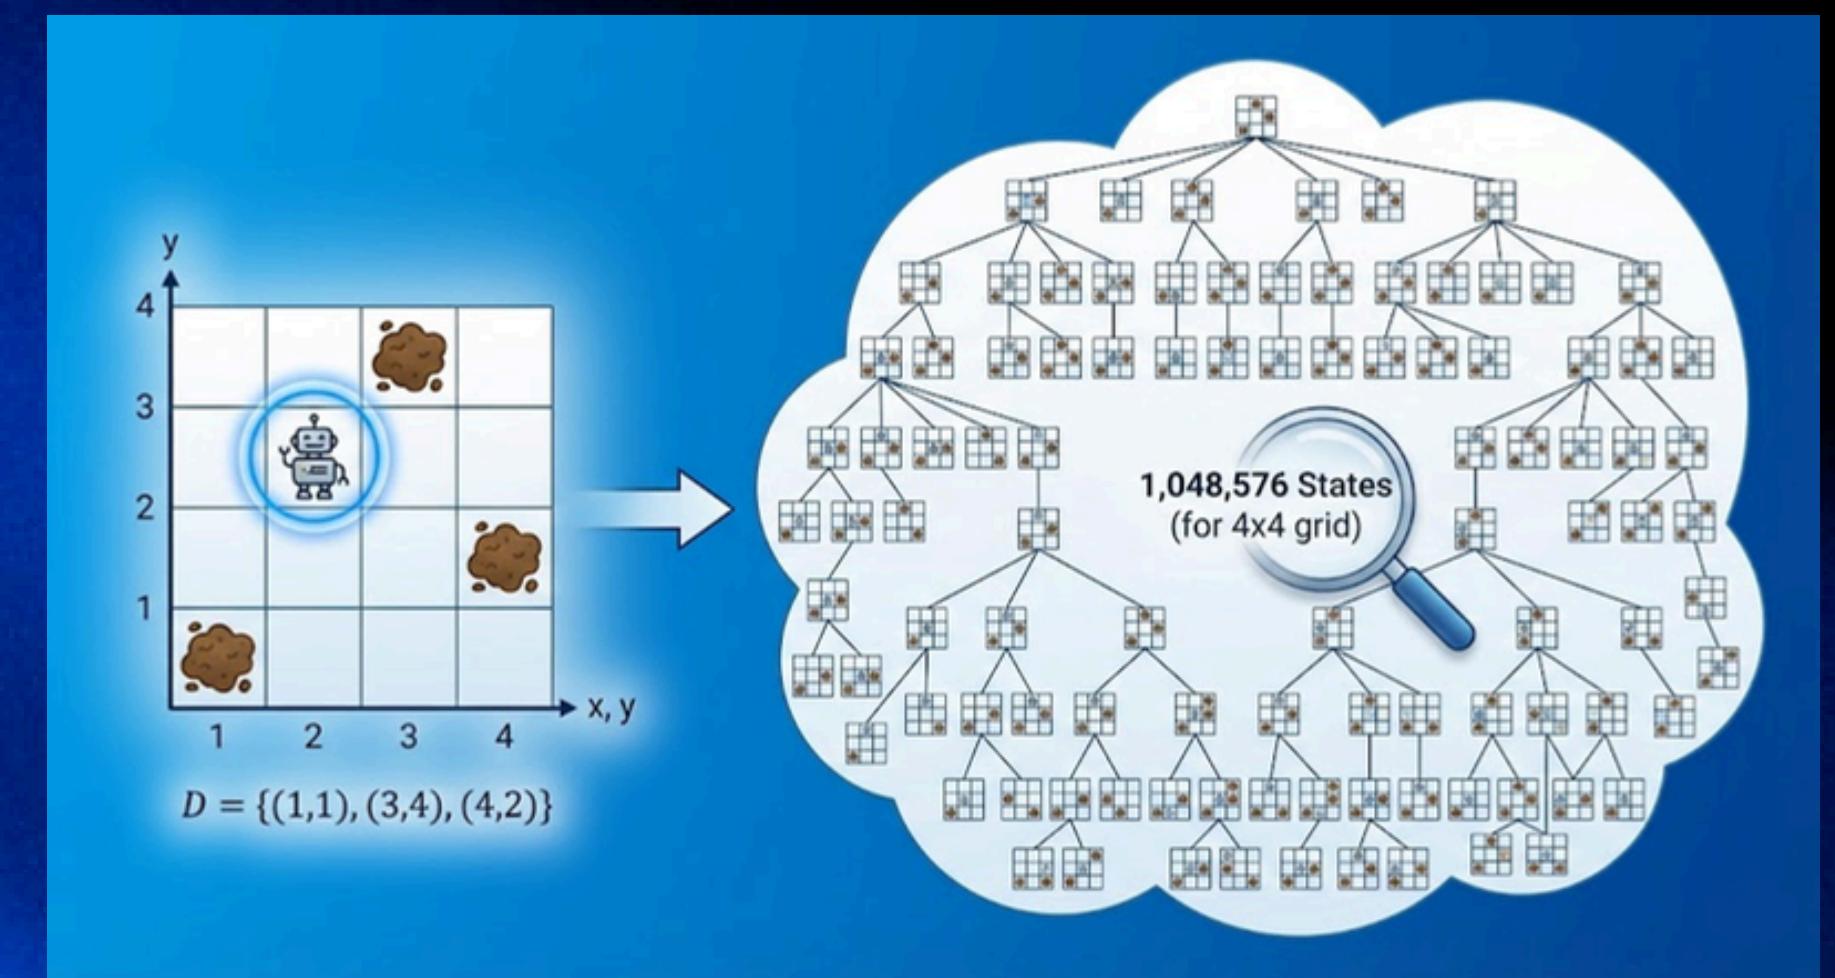


# Formal State Definition

**Formula:**  $\text{State} = (x, y, D)$

- $(x, y)$ : Robot coordinates.
- $D$ : A set of coordinates  $\{(x_1, y_1), (x_2, y_2), \dots\}$  representing dirty cells.

Note: The state space size is  $n^2 \times 2^{n^2}$ . For a  $4 \times 4$  grid, this is  $16 \times 2^{16} = 1,048,576$  states!



# Action Set & Goal Test

## Action Set

To interact with the  $n \times n$  grid, the agent is equipped with 5 discrete actions:

- Movement Actions: {Up, Down, Left, Right}
  - **Logic:** Modifies the agent's  $(x, y)$  coordinates by  $\pm 1$  unit.
  - **Boundary Constraint:** If an action leads to a position outside the grid (e.g.,  $x < 0$  or  $x > n-1$ ), the agent stays put in its current cell.
  - **Cost:** Each movement attempt costs 1 unit, regardless of success.
- Cleaning Action: {Suck}
  - **Logic:** Checks if the current  $(x, y)$  location exists in the dirt set  $D$ .
  - **Effect:** If dirty, the cell is removed from  $D$  (becomes clean). If already clean, the state remains unchanged.
  - **Cost:** Costs 1 unit.



# Action Set & Goal Test

## Goal Test

The search algorithm constantly monitors the state to identify the completion of the mission:

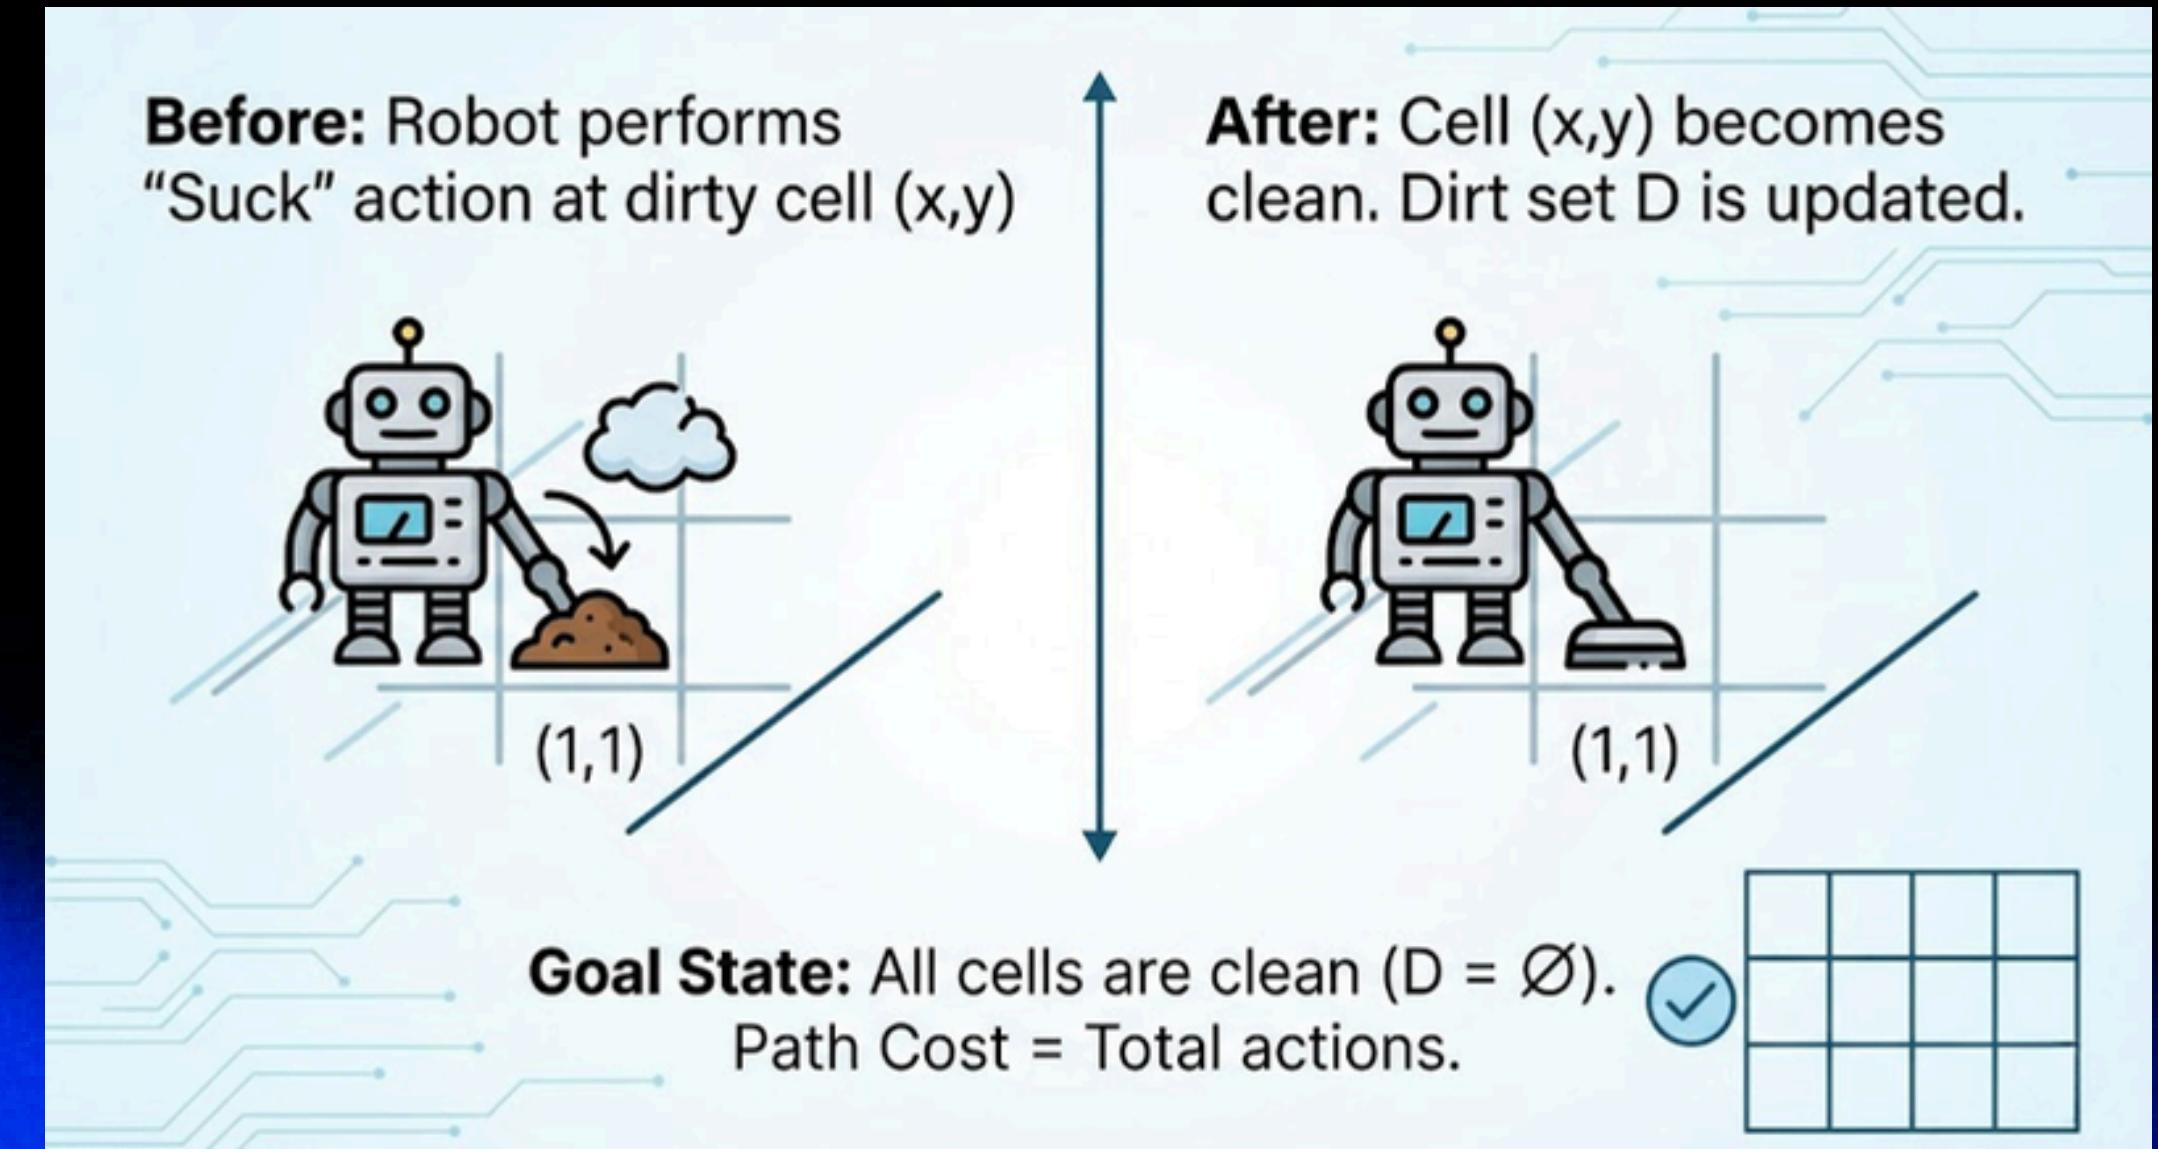
- **Condition:** A state  $s = (x, y, D)$  is identified as a Goal State if and only if:

$$D = \emptyset$$

- **Description:** The set of dirty cells must be empty, meaning every cell in the grid is clean.
- **Termination:** Once this condition is met, the agent stops and returns the optimal sequence of actions.



# Action Set & Goal Test



## Path Cost Function

**Step Cost:**  $c(s, a, s') = 1$ .

**Total Path Cost:** The sum of all individual step costs from the initial state to the goal.

**Objective:** The agent's primary directive is to find a path where the total cost is minimized.

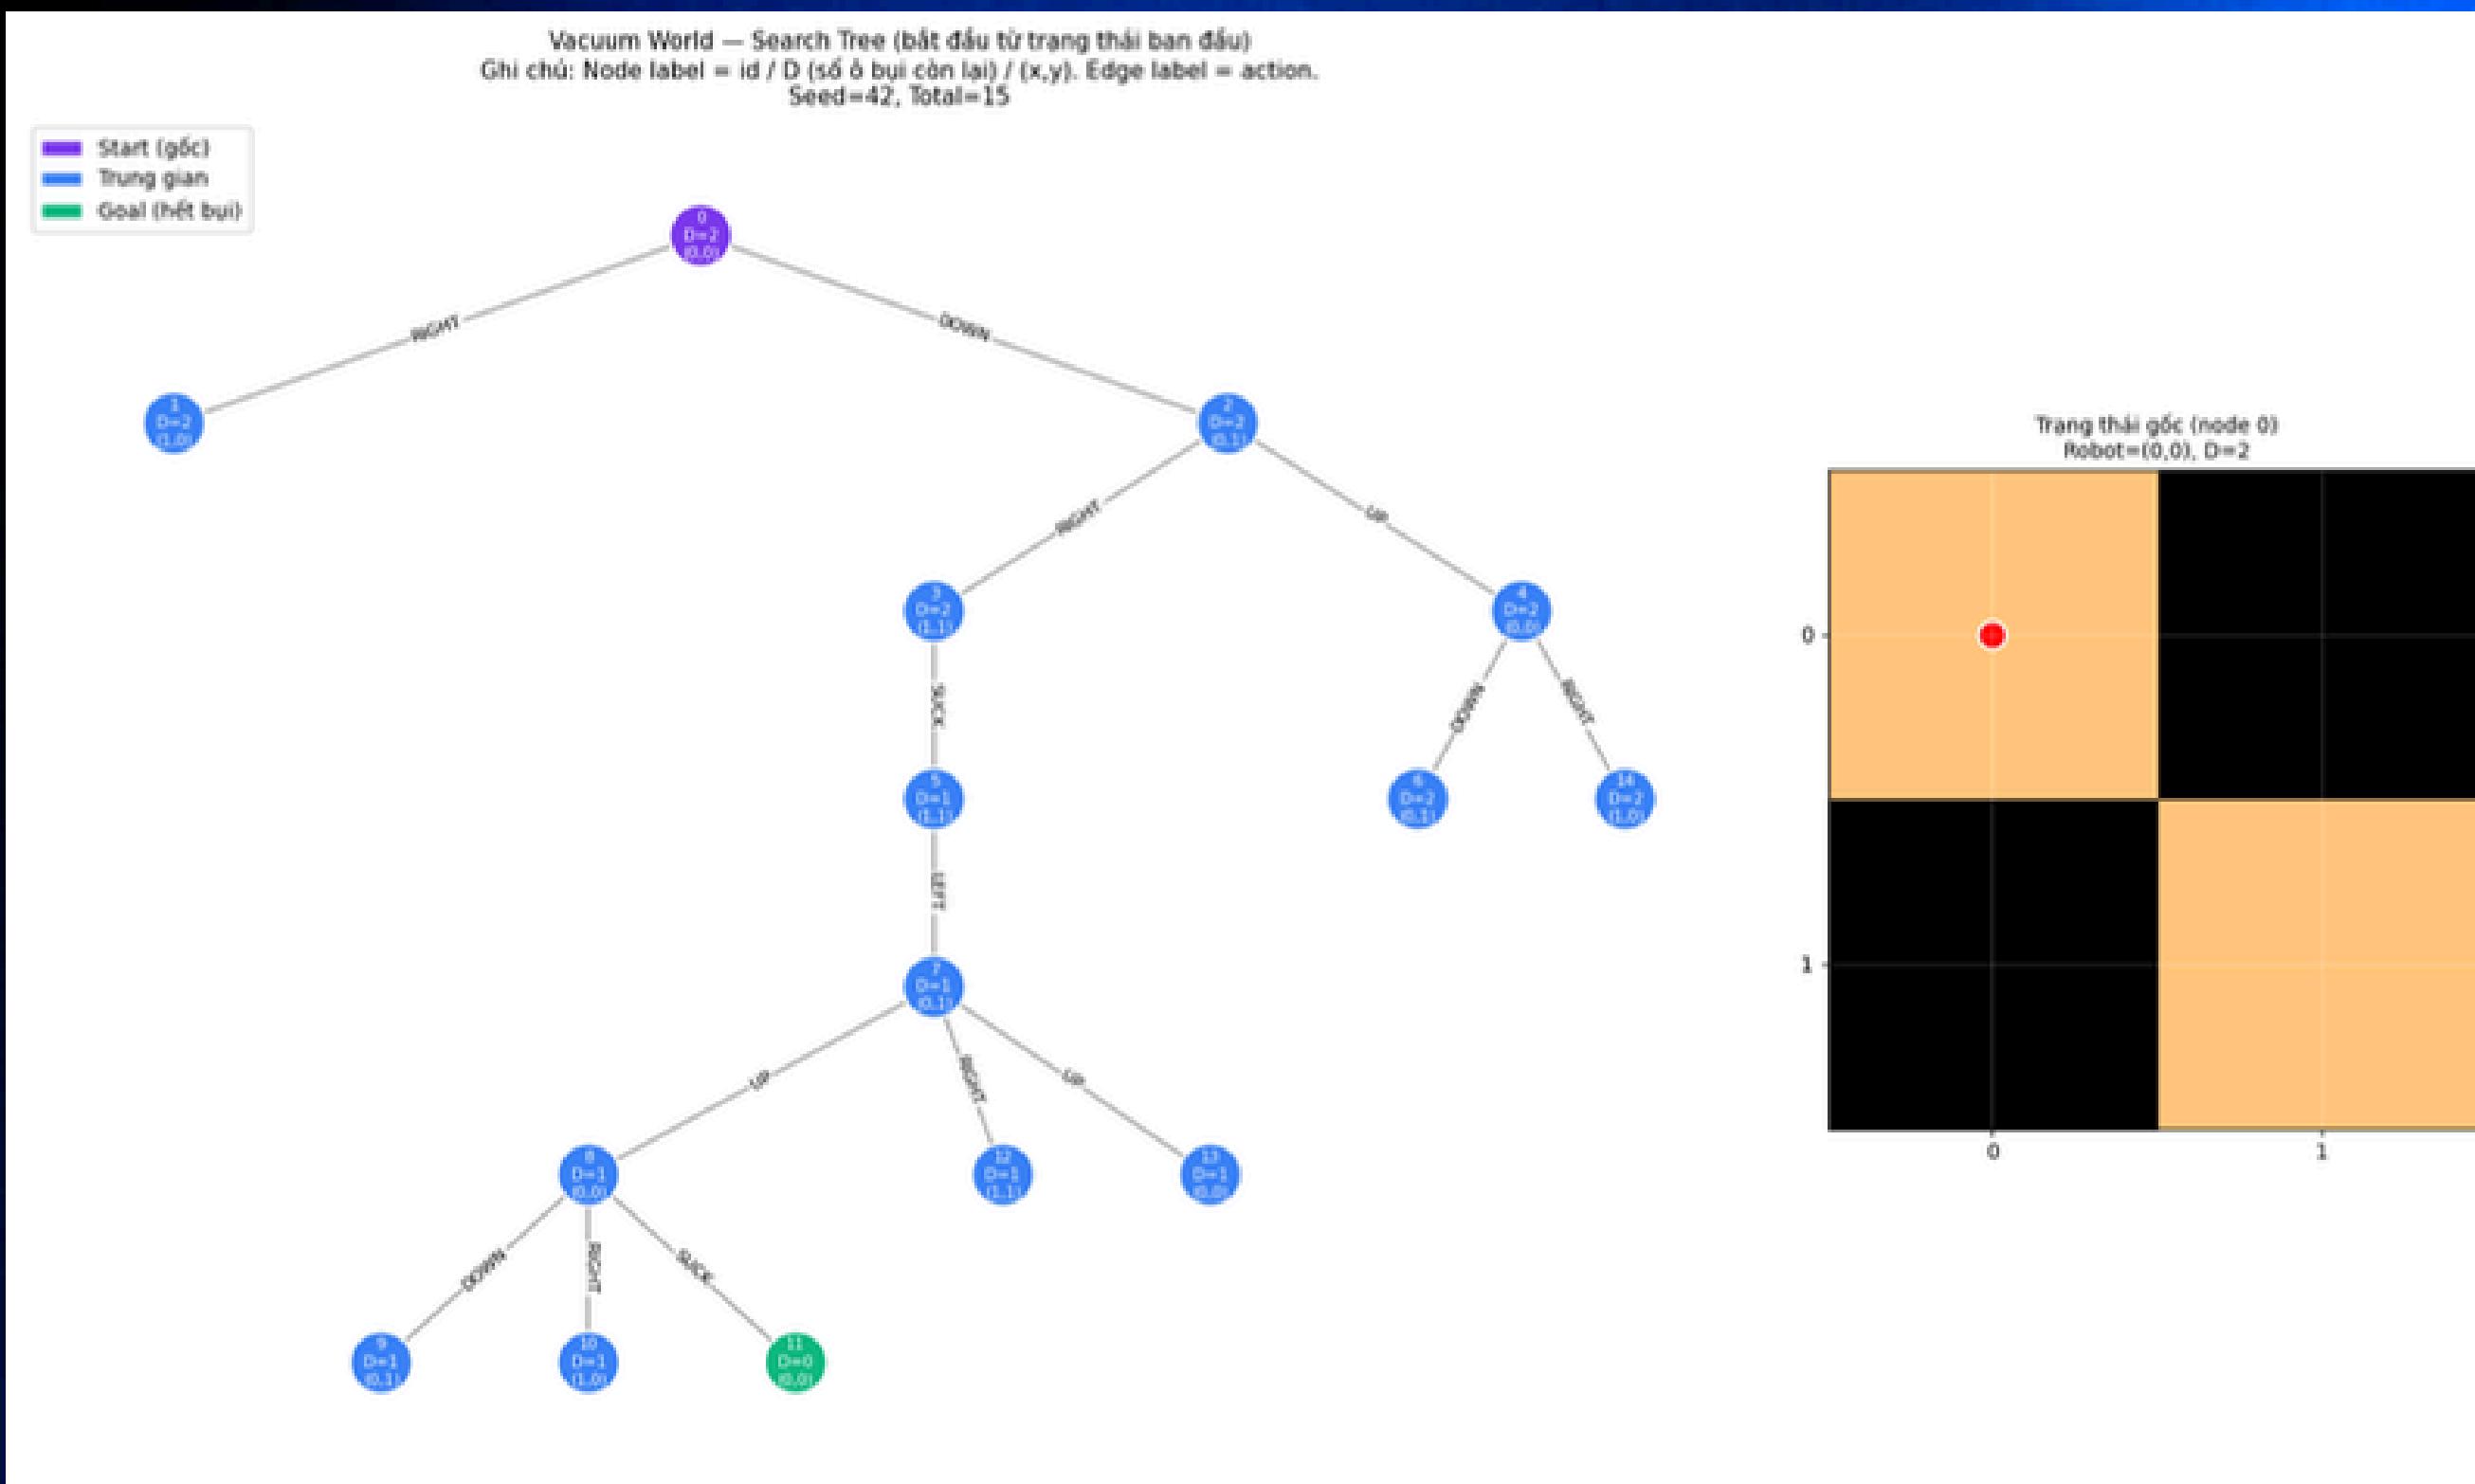


# State Transitions Function

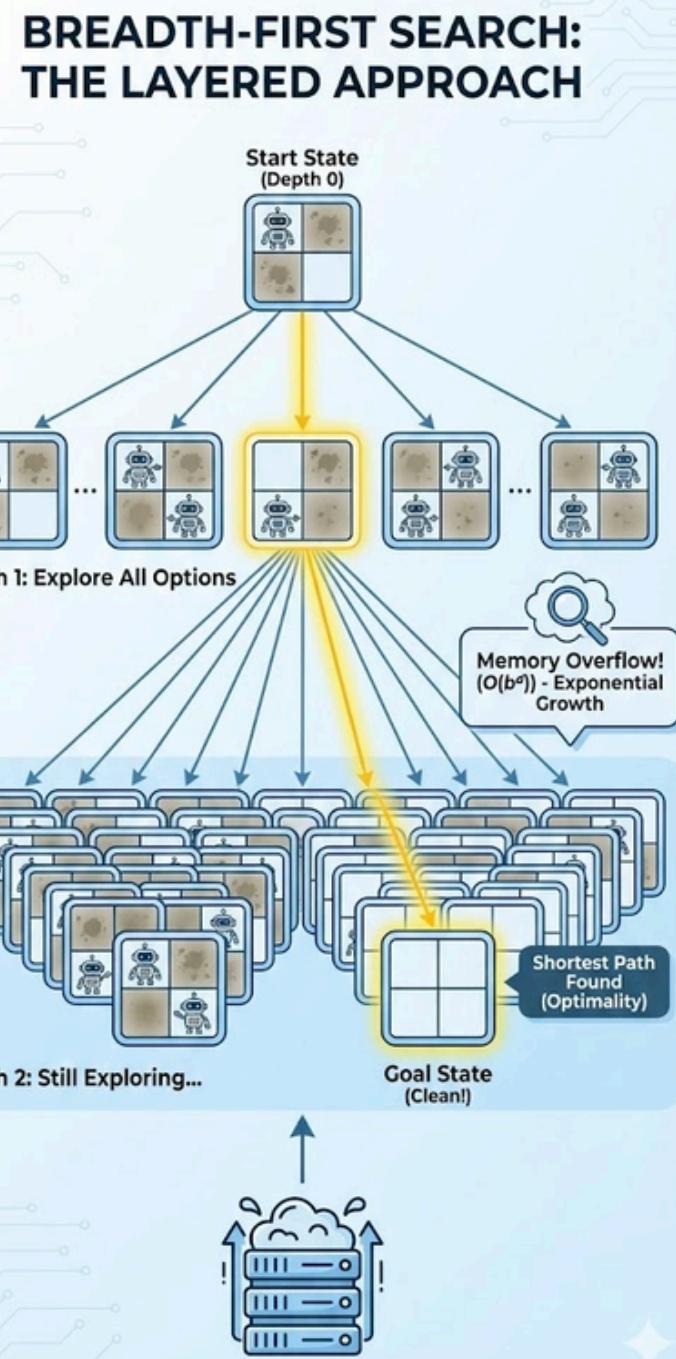
No.	Current state (s) (Robot position; Dirt list)	Action (a)	Next state(s') (Robot position; Dirt list)	Explanation (Execution logic)
1	Robot: (2,2); Dirt: {(1,1),(3,3)}	Up	Robot: (1,2); Dirt: {(1,1),(3,3)}	MOVEMENT: Successfully moved onto the empty space.
2	Robot: (1,2); Dirt: {(1,1),(3,3)}	Left	Robot: (1,1); Dirt: {(1,1),(3,3)}	MOVEMENT: Into the dirty square, not yet sucked up dirt.
3	Robot: (1,1); Dirt: {(1,1),(3,3)}	Suck	Robot: (1,1); Dirt: {(3,3)}	Dirt SUCTION: At (1,1) there is Dirt -> Cleaned.
4	Robot: (1,1); Dirt: {(3,3)}	Suck	Robot: (1,1); Dirt: {(3,3)}	Dirt SUCTION: Cell (1,1) is clean -> No change.
5	Robot: (1,1); Dirt: {(3,3)}	Up	Robot: (1,1); Dirt: {(3,3)}	WALL: Top Border -> Stay still.
6	Robot: (1,1); Dirt: {(3,3)}	Left	Robot: (1,1); Dirt: {(3,3)}	WALL: Left Border -> Stay still.



# Search Tree Visualization



# Search Strategies



## Concept:

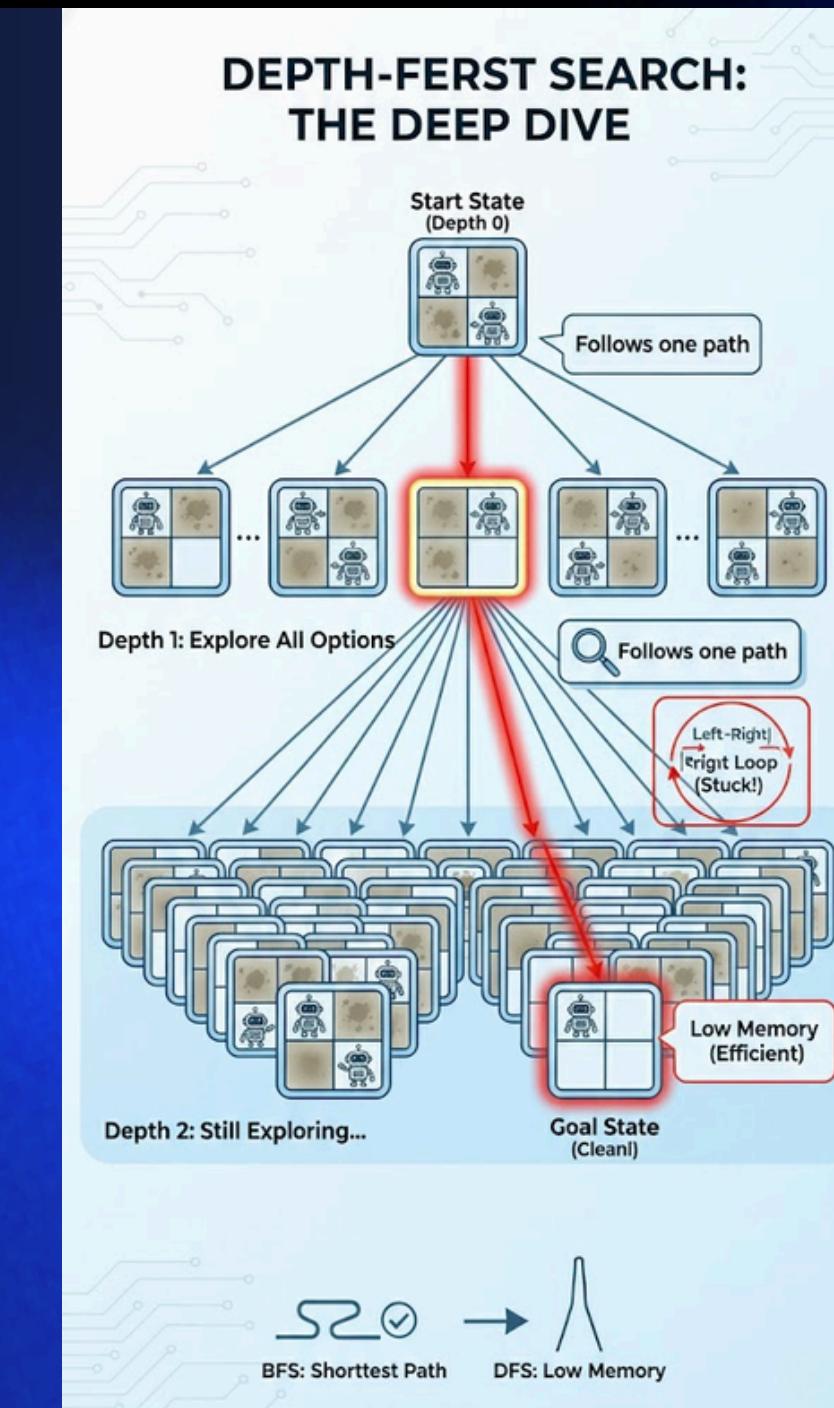
Explores all nodes at depth  $d$  before  $d+1$ .

## Pros:

Guaranteed to find the shortest path (Optimality).

## Cons:

Exponential memory growth ( $O(b^d)$ ). In Vacuum World, "Depth" is the number of moves to clean everything.



## Concept:

Follows one path until a leaf or goal is reached.

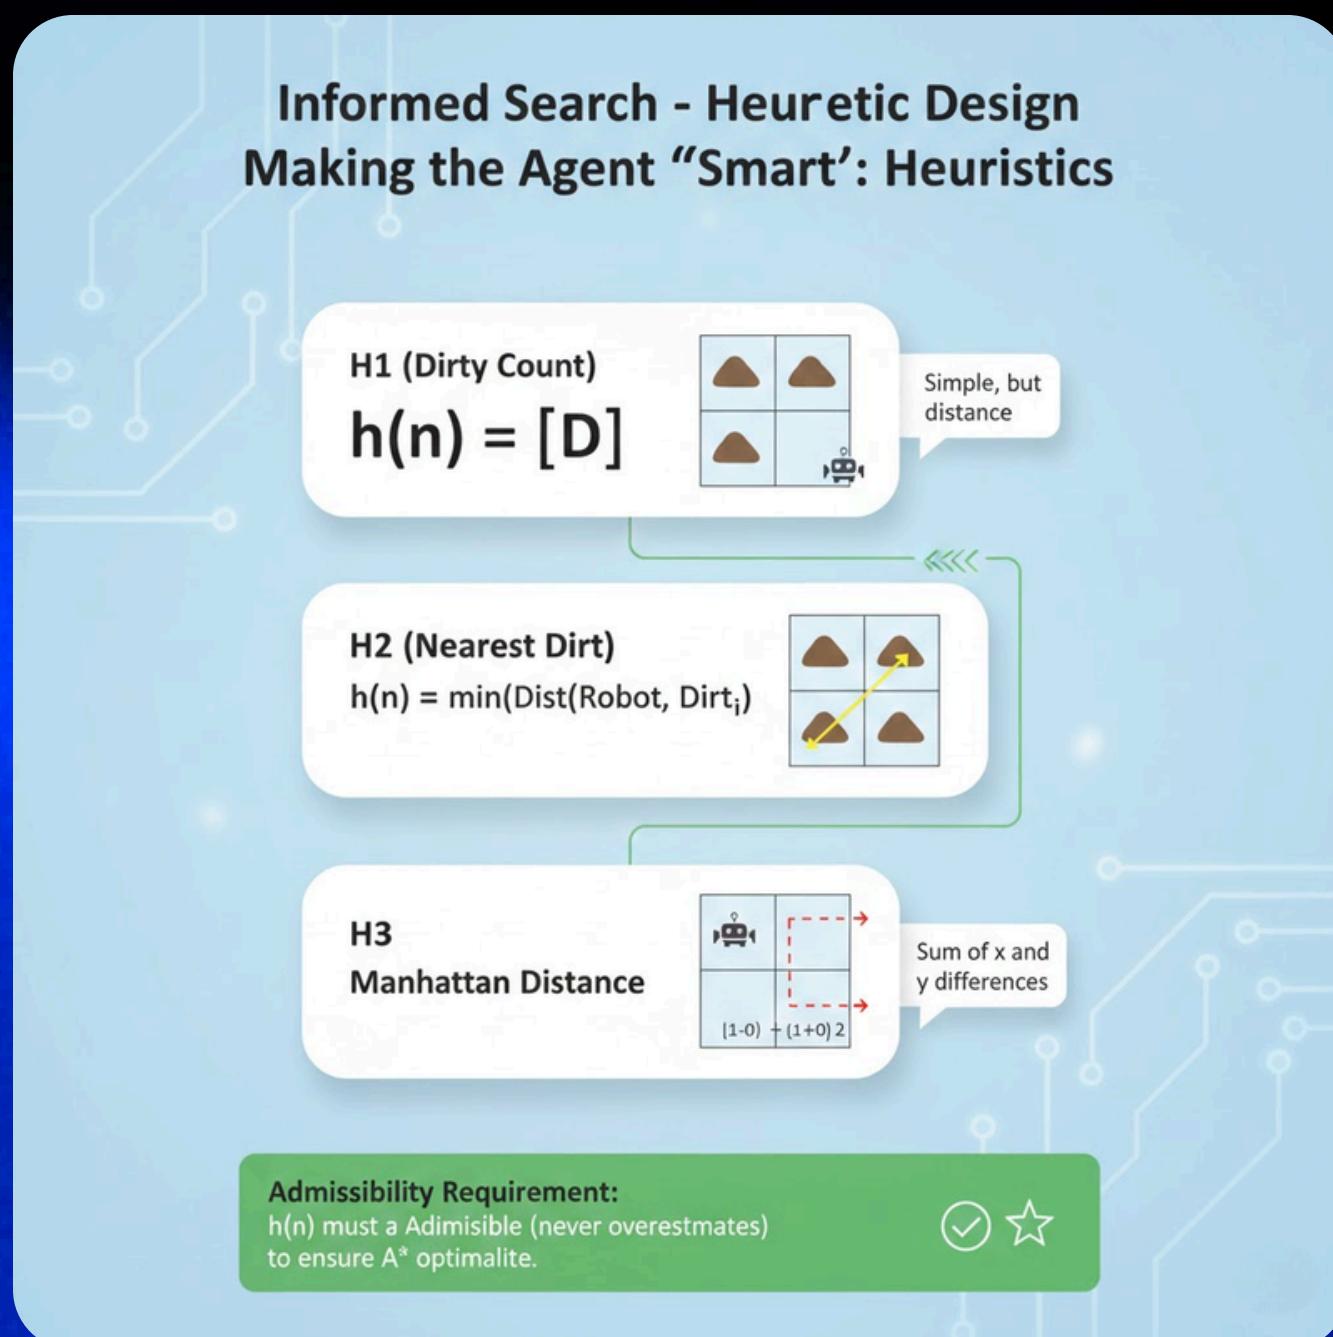
## Pros:

Can get stuck in "Movement Loops" (Left-Right-Left) without an Explored Set.

## Cons:

BFS is better for "Shortest Path", DFS is better for "Low Memory".

# Making the Agent "Smart": Heuristics



Concept:

Follows one path until a leaf or goal is reached.

Pros:

Can get stuck in "Movement Loops" (Left-Right-Left) without an Explored Set.

Cons:

BFS is better for "Shortest Path", DFS is better for "Low Memory".

# Heuristic Design

What is a Heuristic ( $h(n)$ )?

- **Definition:** A function that estimates the cheapest cost from the current state  $n$  to a Goal State.
- **Role:** It acts as a "compass," guiding the search algorithm toward promising states and pruning irrelevant paths.
- **Admissibility Condition:** For A\* to guarantee an optimal solution, the heuristic must never overestimate the true cost ( $h(n) \leq h^*(n)$ ).



# Algorithm Performance

**Strategy 1: Dirt-Count Heuristic ( $h_1$ )**

**Formula:**

$$h_1(n) = |D|$$

**Strategy 2: Manhattan Distance ( $h_2$ )**

**Formula:**

$$h_{\text{Manhattan}}(A, B) = |x_A - x_B| + |y_A - y_B|$$

**Application in Vacuum World:**

$$h_2(n) = \min(\text{Manhattan}(\text{Robot}, d))$$

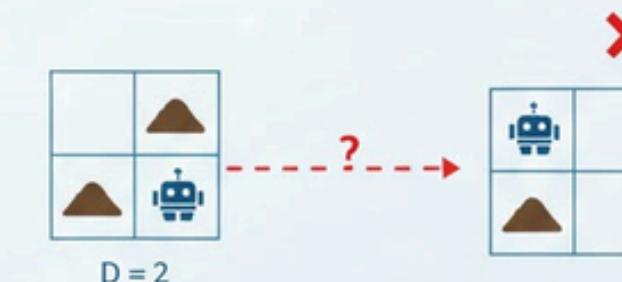
## Informed Search: Heuristic Strategies

**Strategy 1: Dirt-Count Heuristic  $h_1$**

$$h_1(n) = [D]$$



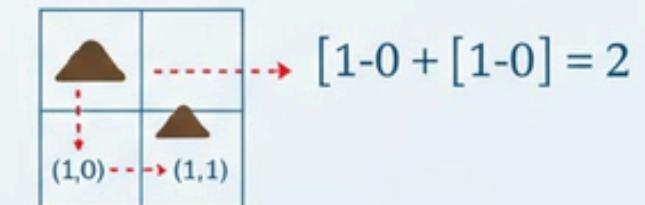
- ✓ **Pros:** Fast to compute ( $O(1)$ )
- ✓ **Cons:** Ignores distance (Weak)



**Strategy 2: Manhattan Distance ( $h_2$ )**

$$h_2(n) = \min \text{ d en } D$$

(Manhattan(Robot, d))



- **Logic:** Grid movement

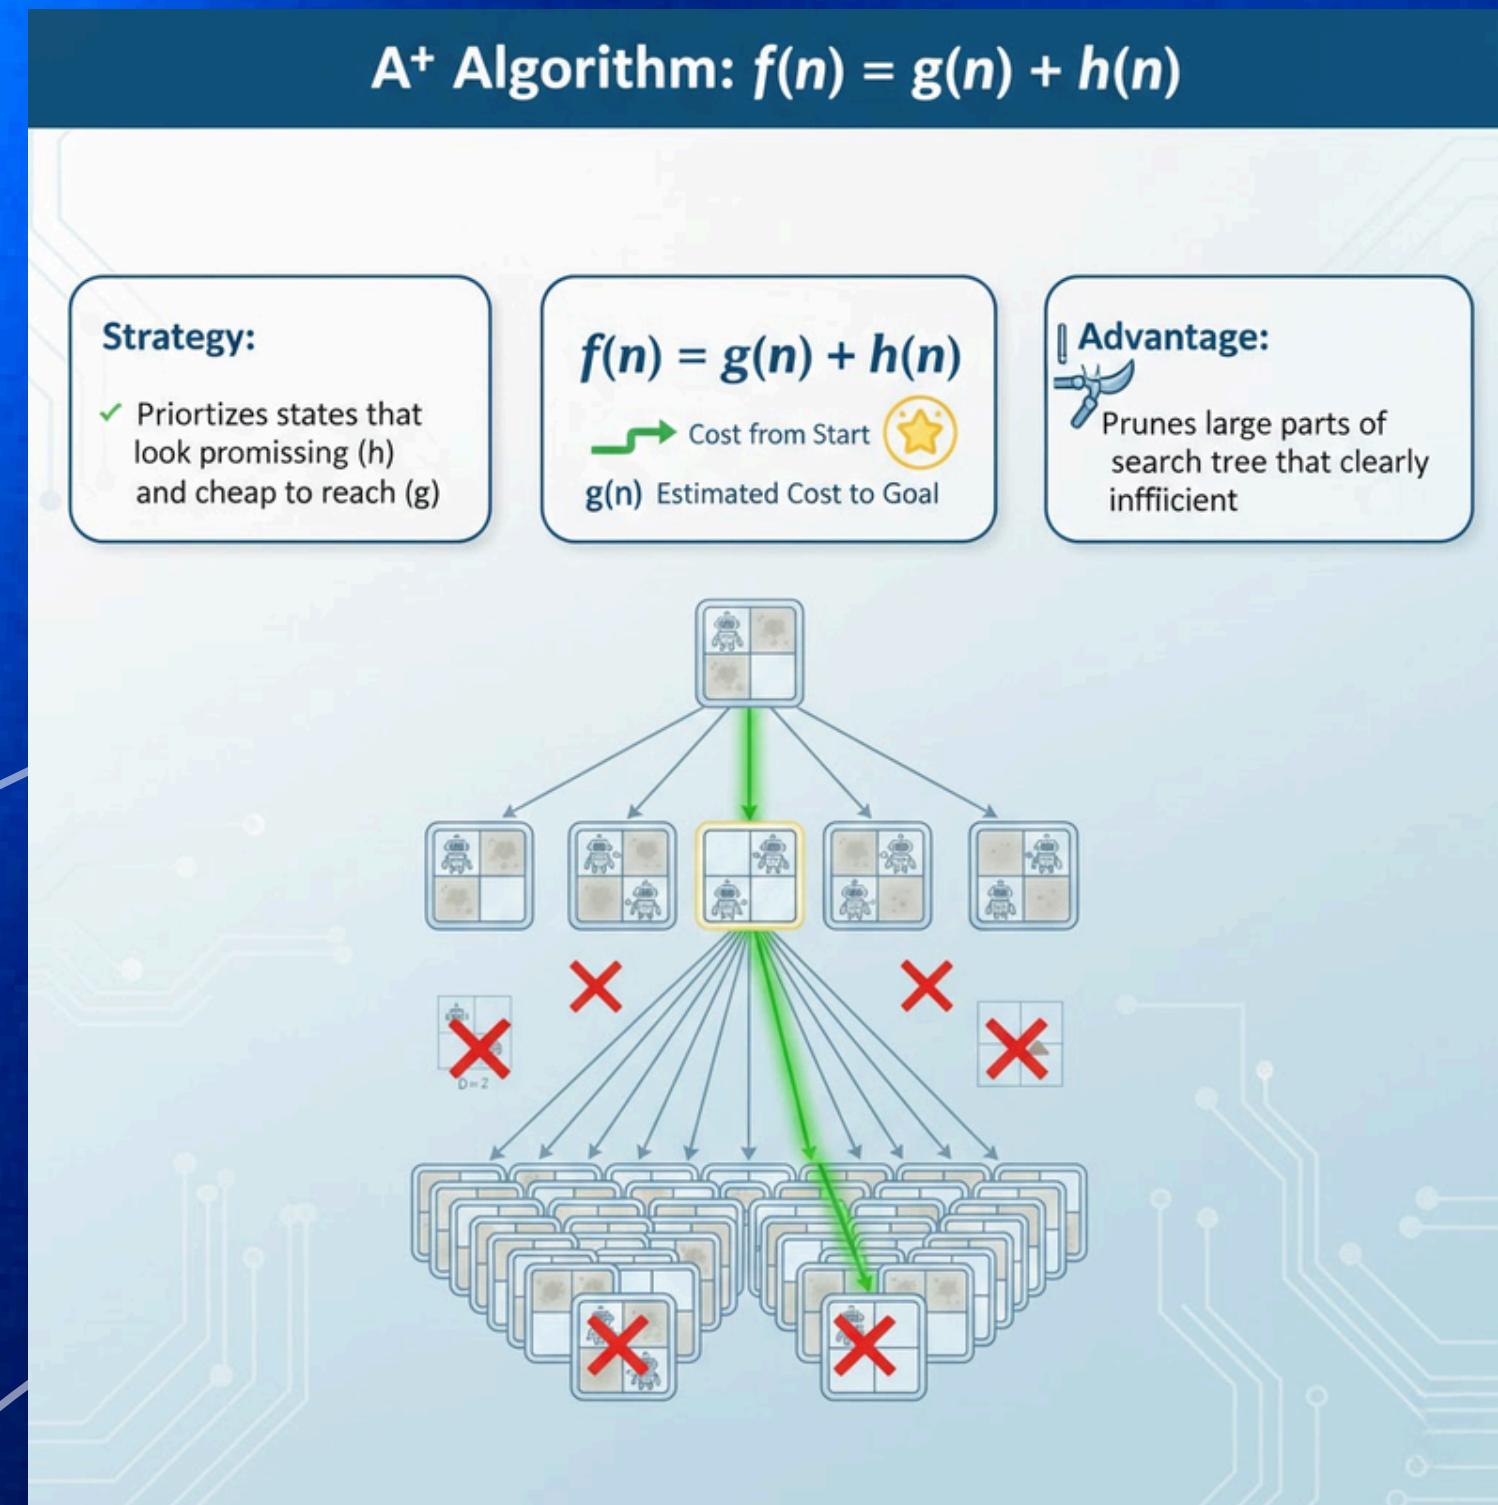


- **Application:** Nearest Dirty Cell

Goal: Find shortest path using smart estimates.

# A Search

## Efficiency Meets Optimality\*



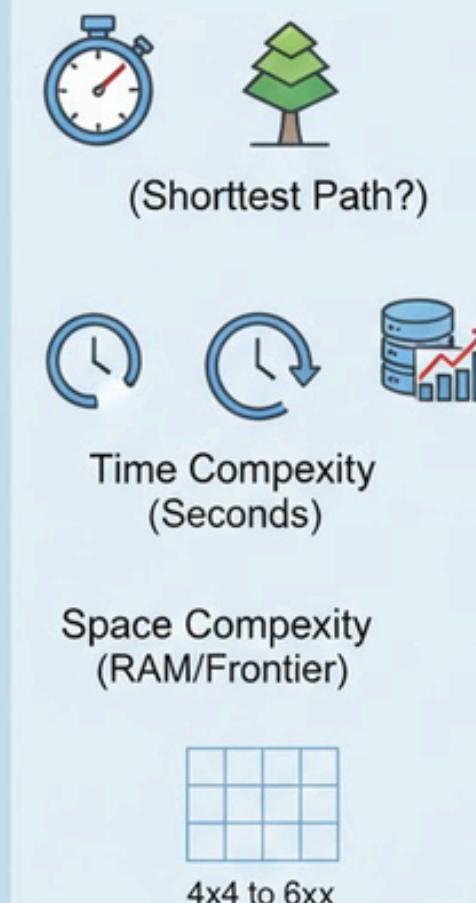
# Algorithm Performance

To judge our agents, we measured three key metrics on varying grid sizes (4x4 to 6x6):

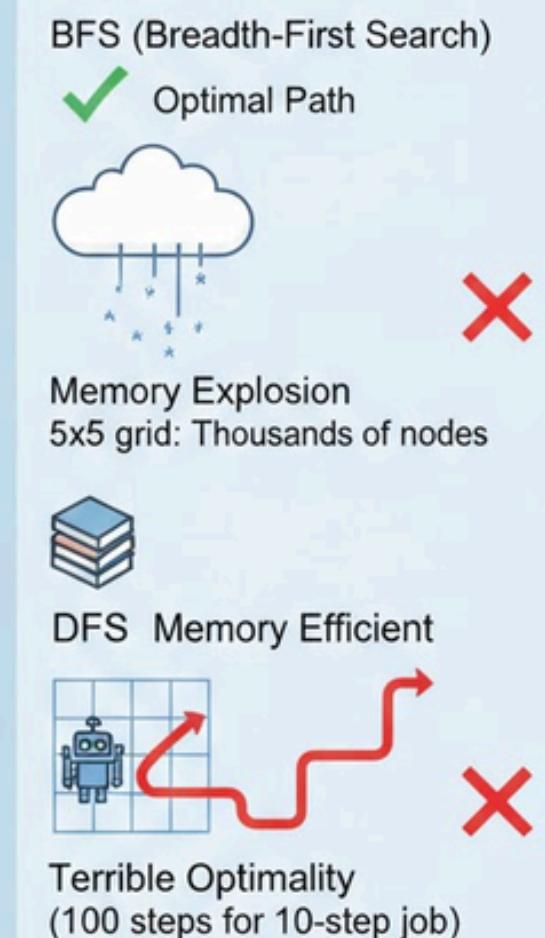
- **Optimality:** Does the agent find the path with the minimum number of steps?
- **Time Complexity:** How long does it take to run? (Measured in seconds/nodes expanded).
- **Space Complexity:** How much memory (RAM) does the "Frontier" queue consume?

## Evaluation Metrics & Algorithm Comparison

### 1. Evaluation Metrics



### 2. Uniformed Search



### 3. Informed Search



Goal: Find shortest path efficiently.

# "Graph Search"

VS

# "Constructive Heuristics"

## Scaling Up: Tacking the State Space Explosion

### 1. The "State Space Explosion"

$$n^2 \times 2^{n^2}$$

✓ Manageable



10x4  
Exceeds Memory



Computational Wall  
(BFS, DFS, A\* fail)

### 2. Two Approaches to Problem Solving

Approach A: Graph Search  
(A\*, BFS)



Verdict: ✓ Perfect for small tasks

Approach B: Constructive Heuristic (Greedy NN)  
(Greedy NN)

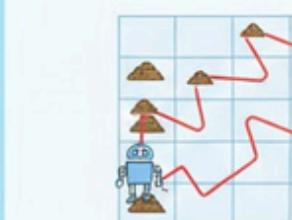


Verdict: ✓ Perfect for small tasks  
Essential for large tasks

### 3. The Solution: Greedy Neighbor (Greedy NN)

$$\text{⌚ } O(k^2 \cdot n^2)$$

Millions for 20x20



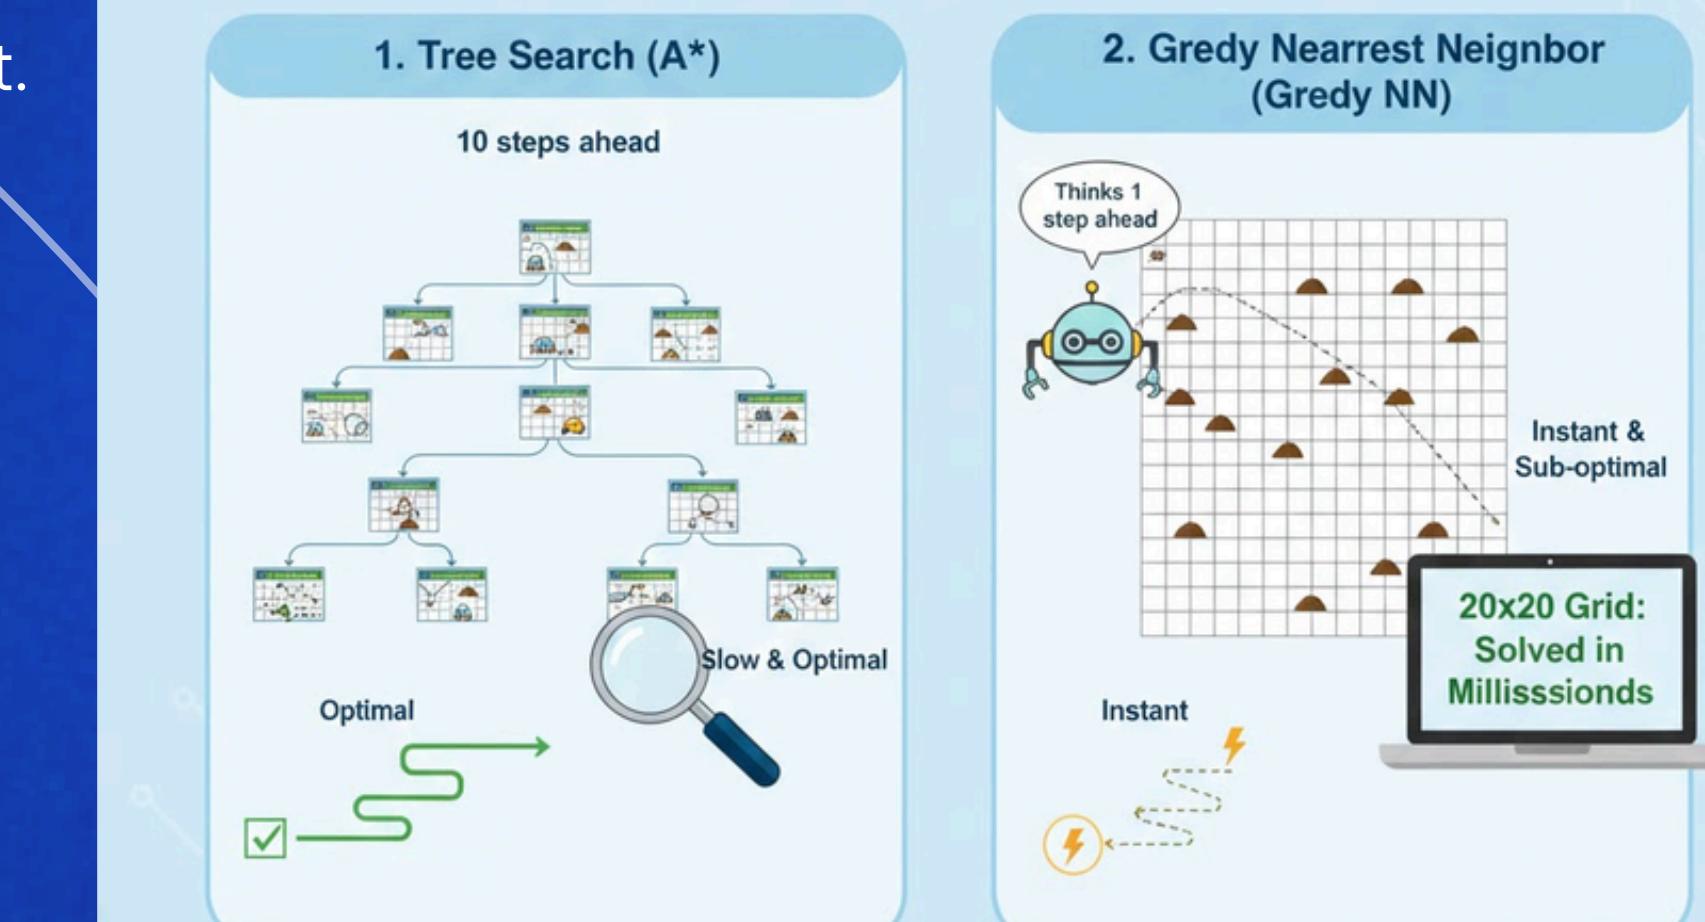
Sub-optimal path  
✓ Gets job done instantly

Goal: Solve large-scale problems quickly

# Constructive Heuristic vs. Search

- **Greedy NN:** Don't search a tree; just move to the closest dirt.
- **Comparison:**
- Search:** Thinks 10 steps ahead (Slow, Optimal).
- Greedy NN:** Thinks 1 step ahead (Instant, Sub-optimal).
- **Result:** Greedy NN is the only viable option for 20 x 20 grids.

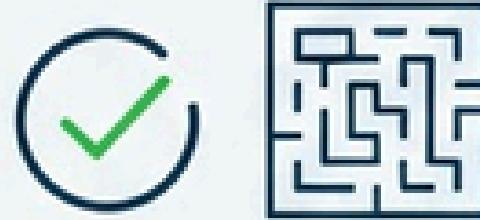
## Greedy Nearest Neighbor (Greedy NN) in Action



DEMO

# Experimental Results & Algorithm Comparison

## Completeness



Guaranteed to find  
find a solution if  
one exists?

## Optimality



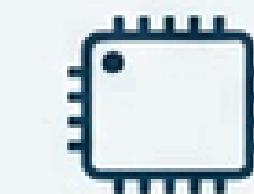
Finds the Shortest  
Shortest Path  
(minimum moves?)

## Time Complexity



$\$O(b^d)$  vs.  $\$O(n^k)$   
How fast as  
problem grows?

## Space Complexity



Frontier &  
Explored

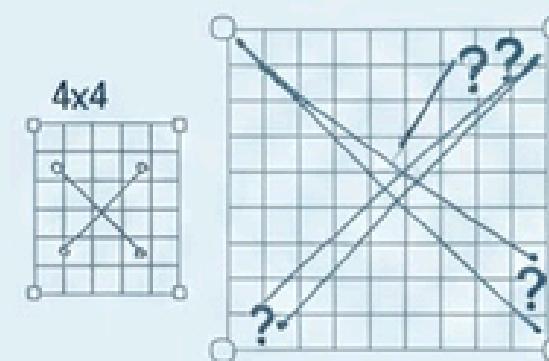
How much RAM  
consumed?

PERFORMANCE  
Metrics

# Comprehensive Comparison Table

Search Category	Algorithm	Optimality	Completeness	Time/Space	Scalability
Uninformed	BFS	Yes	Yes	$O(b^d)$	Very Low (max 5x5)
Uninformed	DFS	No	Yes*	$O(b^m)$	Moderate
Informed	A*	Yes	Yes	$O(b^d)$	Low (max 7x7)
Informed	Greedy Best-First	No	No (Loops)	$O(b^d)$	Low
Constructive	Greedy NN	No	Yes	<i>Polynomial</i>	High (100x100+)

## 1. Graph Search (BFS, A\*)

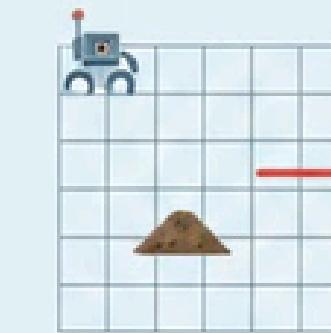


Computational Wall



Algorithms Crash

## 2. Constructive Heuristic (Gredy NN)



Acts: Moves to nearest goal immediately



Result: Extremely fast (<100ms) on 20x20, path 50-100% longer

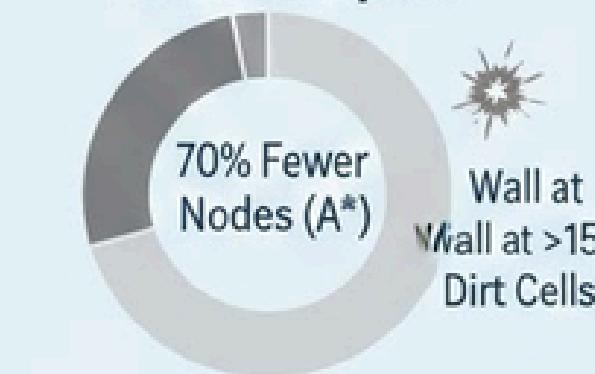
## 3. Key Findings from Our Experiments

Memory Bottleneck



Frontier Queue (BFS)

Heuristic Impact



70% Fewer Nodes (A\*)

Wall at >15 Dirt Cells

The Trade-off



A\* (Min Energy)



(Real-time)

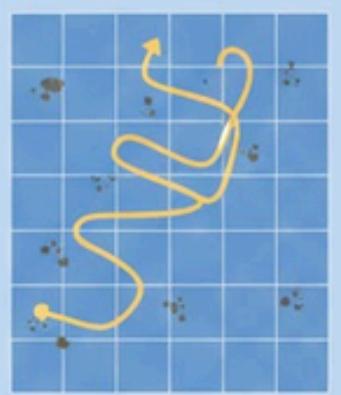
## Algorithm Suitability & Scalability

### 1. Suitability & Scalability Trade-off



A\* Search  
(Gold Standard)

Small-to-Medium Grids:  
Optimal, Efficient



Greedy NN  
(Practical)

Large-Scale Grids:  
Fast, Sub-optimal

Optimality

A\*  
LOW

Scalability

MG  
HIGH

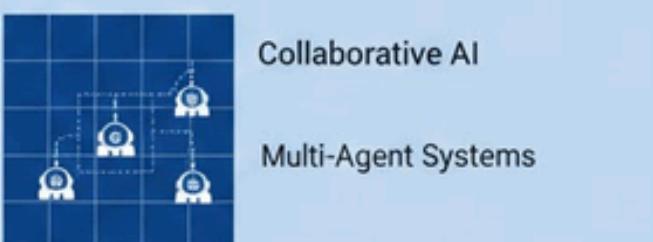
Bridging the gap between simulation and reality

### 2. Future Research Directions



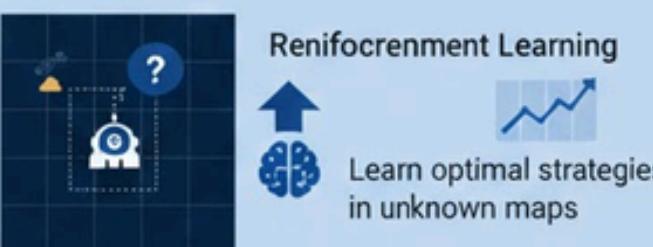
Dynamic Environments

Dirt re-appears over time



Collaborative AI

Multi-Agent Systems



Reinforcement Learning

Learn optimal strategies in  
unknown maps

## Algorithm Suitability & Scalability

The Optimality-Scalability Trade-off: Our analysis confirms that A Search\* is the "Gold Standard" for small-to-medium grids, offering the perfect balance of path precision and memory efficiency—provided the **Heuristic Quality** (e.g., Manhattan distance) is high. However, for large-scale practical applications, we observed that **Constructive Heuristics** (Greedy NN) are strictly necessary, as they trade theoretical optimality for the execution speed required to handle the exponential growth of state spaces.

## Future Research Directions

- **From Static to Dynamic Intelligence:** The next phase of development aims to bridge the gap between simulation and reality by introducing:
- **Dynamic Environments:** Handling stochastic elements where dirt re-appears over time.
- **Collaborative AI:** Implementing Multi-Agent Systems to clean the grid in parallel.
- **Reinforcement Learning:** allowing the agent to learn optimal strategies in unknown map layouts without prior knowledge.

THANK YOU  
FOR LISTENING!

